

3. Multiple Linear Regression

October 23, 2021

Multiple (Multi-variate) Linear Regression

```
[53]: # Import necessary package
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

0.0.1 Step 1: Load the dataset

```
[54]: # Load the dataset into pandas dataframe
df=pd.read_csv("E:\\MY LECTURES\\DATA SCIENCE\\3.Programs\\dataset\\Advertising.
↪csv")
# Change this location based on the location of dataset in your machine
```

```
[55]: # Display the first five records
df.head()
```

```
[55]:
```

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

Advertising data comprises four features: TV, radio, newspaper, and sales. It explains the budget (in 1000\$) spent on different mass media and the net outcome for every week.

sales for a product (output/dependent/target variable).

advertising budget for TV, radio, and newspaper media (input/independent/target variable).

Planning to perform regression on TV budget (X1), Radio budget (X2), Newspaper budget (X3) as input and sales (Y) as output.

```
[56]: # Dataset shape (number of rows and columns)
df.shape
```

```
[56]: (200, 4)
```

Row <=> record, tuple, instance, sample, observation, object, case, entity Column <=> attribute, variable, field, feature, characteristic, dimension

0.0.2 Step 2: Apply EDA

Univariate analysis

```
[57]: # Statistics summary  
df["TV"].describe()
```

```
[57]: count      200.000000  
mean       147.042500  
std        85.854236  
min         0.700000  
25%        74.375000  
50%       149.750000  
75%       218.825000  
max       296.400000  
Name: TV, dtype: float64
```

```
[58]: df["radio"].describe()
```


```
[58]: count      200.000000  
mean        23.264000  
std         14.846809  
min          0.000000  
25%          9.975000  
50%         22.900000  
75%         36.525000  
max         49.600000  
Name: radio, dtype: float64
```

```
[59]: df["newspaper"].describe()
```

```
[59]: count      200.000000  
mean        30.554000  
std         21.778621  
min          0.300000  
25%         12.750000  
50%         25.750000  
75%         45.100000  
max        114.000000  
Name: newspaper, dtype: float64
```


```
[60]: # Univariate Analysis using Boxplot
sns.set_style(style='white')
df.boxplot(column=['TV'], grid=False,figsize=(8,5))
```

[60]: <AxesSubplot:>




```
[61]: df.boxplot(column=['newspaper'], grid=False,figsize=(8,5))
```

[61]: <AxesSubplot:>




```
[62]: df.boxplot(column=['radio'], grid=False,figsize=(8,5))
```

```
[62]: <AxesSubplot:>
```



```
[63]: df.boxplot(grid = False,figsize=(8,5))
```

```
[63]: <AxesSubplot:>
```



```
[64]: # Distribution plot to find skewness
from pylab import *
sns.set(rc={"figure.figsize": (20, 15)});

subplot(4,4,1)
ax = sns.histplot(df["TV"])

subplot(4,4,2)
ax = sns.histplot(df["TV"], kde=True, stat="density", linewidth=0)

subplot(4,4,3)
ax = sns.histplot(df["radio"])

subplot(4,4,4)
ax = sns.histplot(df["radio"], kde=True, stat="density", linewidth=0)

subplot(4,4,5)
ax = sns.histplot(df["newspaper"])
```

```


subplot(4,4,6)
ax = sns.histplot(df["newspaper"], kde=True, stat="density", linewidth=0)

subplot(4,4,7)
ax = sns.histplot(df["sales"])

subplot(4,4,8)
ax = sns.histplot(df["sales"], kde=True, stat="density", linewidth=0)

plt.show()

```




Bivariate analysis

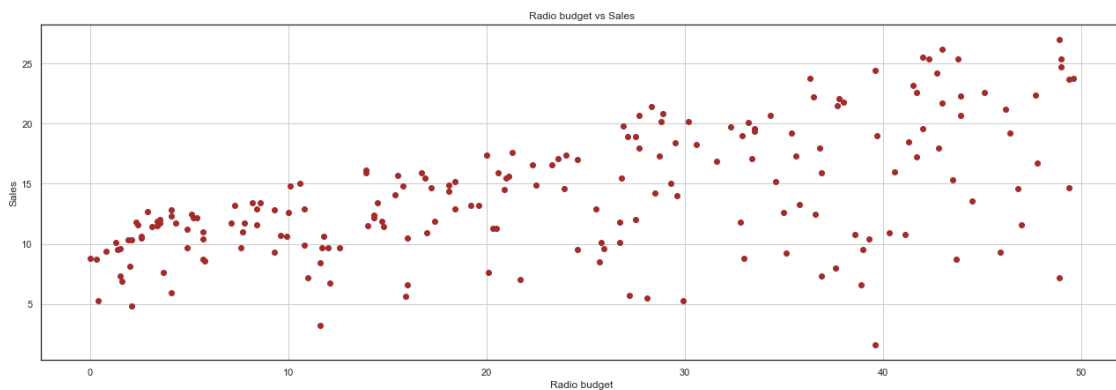
```

[65]: # Scatter plot
sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
plt.scatter(df["TV"],df["sales"],color="brown")
plt.grid(b=None)
plt.xlabel("TV budget")
plt.ylabel("Sales")
plt.title("TV budget vs Sales")
plt.show()

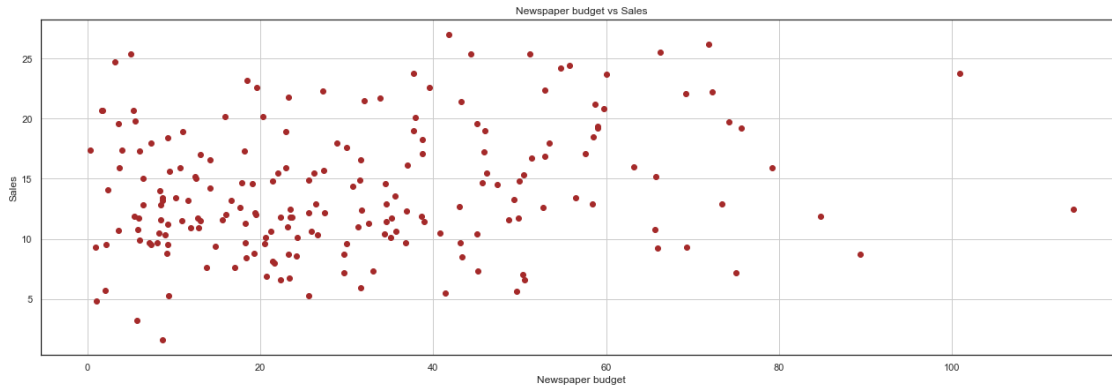
```



```
[66]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
plt.scatter(df["radio"],df["sales"],color="brown")
plt.grid(b=None)
plt.xlabel("Radio budget")
plt.ylabel("Sales")
plt.title("Radio budget vs Sales")
plt.show()
```



```
[67]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
plt.scatter(df["newspaper"],df["sales"],color="brown")
plt.grid(b=None)
plt.xlabel("Newspaper budget")
plt.ylabel("Sales")
plt.title("Newspaper budget vs Sales")
plt.show()
```



```
[68]: # Correlation
df.corr()
# Spearman's rho
# df.corr(method='spearman')
# Kendall's tau
# df.corr(method='kendall')
```

```
[68]:
```

	TV	radio	newspaper	sales
TV	1.000000	0.054809	0.056648	0.782224
radio	0.054809	1.000000	0.354104	0.576223
newspaper	0.056648	0.354104	1.000000	0.228299
sales	0.782224	0.576223	0.228299	1.000000

0.0.3 Step 3. Pre-process and extract the features

```
[69]: # Load TV, and radio, and newspaper X1, X2, X3 as feature vector and sales into
      ↪ Y variable
X = df.iloc[:,0:3].values # TV, radio, newspaper
Y = df.iloc[:,3].values  # Sales
```

```
[70]: X[:10,:]
```

```
[70]: array([[230.1,  37.8,  69.2],
          [ 44.5,  39.3,  45.1],
          [ 17.2,  45.9,  69.3],
          [151.5,  41.3,  58.5],
          [180.8,  10.8,  58.4],
          [  8.7,  48.9,  75. ],
          [ 57.5,  32.8,  23.5],
          [120.2,  19.6,  11.6],
          [  8.6,   2.1,   1. ],
          [199.8,   2.6,  21.2]])
```



```
[71]: Y[:10]
```

```
[71]: array([22.1, 10.4,  9.3, 18.5, 12.9,  7.2, 11.8, 13.2,  4.8, 10.6])
```

input feature independent feature or predictor feature. Here, X1 (TV), X2 (Radio), X3 (Newspaper) are the input features. output feature dependent feature or response feature or target feature. Here, Y (sales) is the output feature.

0.0.4 Step 4. Split the data for training and testing

```
[72]: # Splitting dataset into training and testing set
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.2,
↳ random_state = 0)
```

```
[73]: x_train[:5,:]
```

```
[73]: array([[ 36.9,  38.6,  65.6],
           [ 31.5,  24.6,   2.2],
           [142.9,  29.3,  12.6],
           [209.6,  20.6,  10.7],
           [215.4,  23.6,  57.6]])
```

```
[74]: x_test[:5,:]
```

```
[74]: array([[ 69.2,  20.5,  18.3],
           [ 50. ,  11.6,  18.4],
           [ 90.4,   0.3,  23.2],
           [289.7,  42.3,  51.2],
           [170.2,   7.8,  35.2]])
```

```
[75]: y_train[:5]
```

```
[75]: array([10.8,  9.5, 15. , 15.9, 17.1])
```

```
[76]: y_test[:5]
```

```
[76]: array([11.3,  8.4,  8.7, 25.4, 11.7])
```

0.0.5 Step 5: Training phase (bulding the model)

```
[77]: # Fitting line on multiple dimensions on the training set
from sklearn import linear_model
model = linear_model.LinearRegression(normalize=True)
model.fit(x_train, y_train)
```

```
[78]: m1, m2, m3 = model.coef_  
  
[79]: b = model.intercept_  
  
[80]: print("The multi-linear model is :")  
print('\t\t\t\t Y = m1 X1 + m2 X2 + m3 X3 + b \n')  
print('\t\t\t\t Y = ',round(m1,3), 'X1 + ',round(m2,3), 'X2 + ',round(m3,3), 'X3 +  
↵',round(b,3))
```

$$Y = m_1 X_1 + m_2 X_2 + m_3 X_3 + b$$

```
[81]: # Predicting the output for training input
      y_train_pred = model.predict(x_train)
```


Plotting observed sale (x) and predicted sale (y) for training set

```
[82]: x = np.arange(len(y_train_pred))
fig = plt.figure(figsize=(22,7))
plt.plot(x,y_train,"g-+",label="Sales_observed")
plt.plot(x,y_train_pred,"r-*",label="Sales_predicted")
plt.grid(b=None)
plt.xlabel("Record number")
plt.ylabel("Sales")
plt.title("Multi-variate prediction")
plt.legend()
plt.show()
```



Residual (Error) plot If the model has done good predictions, then the datapoints must be near around to horizontal line.

```
[83]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
residuals = y_train-y_train_pred
zeros = y_train-y_train
plt.scatter(y_train_pred,residuals,color="brown")
plt.grid(b=None)
plt.plot(y_train_pred,zeros,"g")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residual plot")
plt.show()
```



0.0.6 Different error calculations to asses the model for training set

1. Sum of Squared Error (SSE)

$$SSE(m,b) = \sum_{i=1}^n (y_i - \hat{y})^2 = \sum_{i=1}^n (y_i - (m * x_i + b))^2 \quad (1)$$

```
[84]: sum = 0
n = len(x_train)
for i in range (0,n):
    diff = y_train[i] - y_train_pred[i]
    squ_diff = diff**2
    sum = sum + squ_diff
Train_SSE = np.round(sum,2)
```

```
print("Sum of Squared Error (SSE) :",Train_SSE)
```

Sum of Squared Error (SSE) : 385.09

2. Mean Squared Error (MSE)

$$MSE(m, b) = \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n} = \frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n} \quad (2)$$

```
[85]: Train_MSE = np.round(Train_SSE/n,2)
print("Mean Squared Error (MSE) :",Train_MSE)
```

Mean Squared Error (MSE) : 2.41

3. Root Mean Squared Error (RMSE)

$$RMSE(m, b) = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n}} = \sqrt{\frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n}} \quad (3)$$

```
[86]: Train_RMSE = np.round(np.sqrt(Train_MSE),2)
print("Root Mean Squared Error (RMSE) :",Train_RMSE)
```

Root Mean Squared Error (RMSE) : 1.55

4. Mean Absolute Error (MAE)

$$MAE(m, b) = \frac{\sum_{i=1}^n |(y_i - \hat{y})|}{n} \quad (4)$$

```
[87]: sum = 0
n = len(x_train)
for i in range (0,n):
    diff = y_train[i] - y_train_pred[i]
    sum = sum + np.abs(diff)
Train_MAE = np.round(sum/n,2)
print("Mean Absolute Error (MAE) :",Train_MAE)
```

Mean Absolute Error (MAE) : 1.21

5. Mean Absolute Percentage Error (MAPE)

$$MAPE(m, b) = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - \hat{y})}{y_i} \right| = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - (m * x_i + b))}{y_i} \right| \quad (5)$$

```
[88]: sum = 0
      n = len(x_train)
      for i in range (0,n):
          diff = (y_train[i] - y_train_pred[i])/y_train[i]
          sum = sum + np.abs(diff)
      Train_MAPE = np.round(sum/n*100,2)
      print("Mean Absolute Percentage Error (MAPE) :",Train_MAPE)
```

Mean Absolute Percentage Error (MAPE) : 11.43

0.0.7 Calculating R-Squared value (goodness of model) using SSE

$$R^2 = 1 - \frac{SSE}{SST} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (6)$$

```
[89]: from sklearn.metrics import r2_score
      out = r2_score(y_train,y_train_pred)
      Train_RS = np.round(out,2)*100
      print("R-Squared value (goodness of model) for training set :",Train_RS,"%")
```

R-Squared value (goodness of model) for training set : 91.0 %

0.0.8 Calculating Adjusted R-Squared value (goodness of model) using SSE

$$R^2 = 1 - (1 - R^2) \frac{(n - 1)}{n - p - 1} \quad (7)$$

```
[90]: out = 1 - (1-Train_RS)*(len(y_train)-1)/(len(y_train)-X.shape[1]-1)
      Train_Adj_RS = round(out,2)
      print("Adjusted R-Squared value (goodness of model) for training set :
      ↪",Train_Adj_RS,"%")
```

Adjusted R-Squared value (goodness of model) for training set : 92.73 %

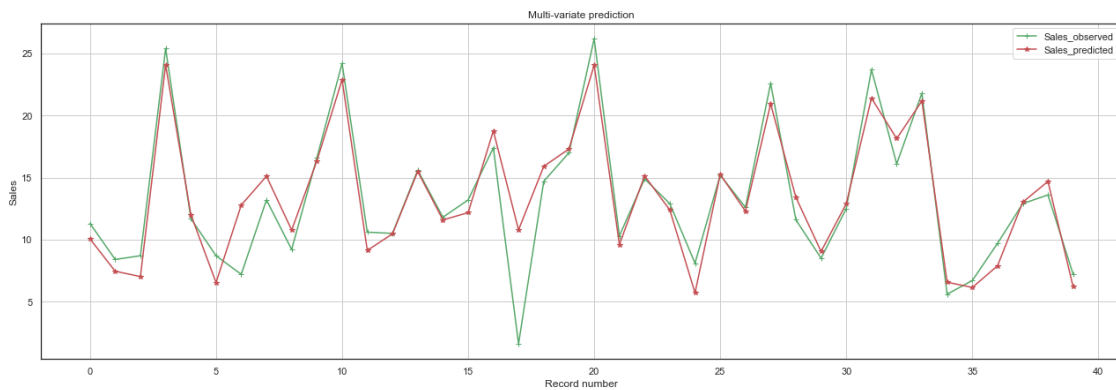
0.0.9 Step 6: Testing phase

```
[91]: # Predicting values for test input set
      y_test_pred = model.predict(x_test)
```

Visualizing the model It involves 4 dimensions, so imagine yourself


Plotting observed sale (x) and predicted sale (y) for test set

```
[92]: x = np.arange(len(y_test_pred))
fig = plt.figure(figsize=(22,7))
plt.plot(x,y_test,"g--",label="Sales_observed")
plt.plot(x,y_test_pred,"r-*",label="Sales_predicted")
plt.grid(b=None)
plt.xlabel("Record number")
plt.ylabel("Sales")
plt.title("Multi-variate prediction")
plt.legend()
plt.show()
```



Residual (Error) plot If the model has done good predictions, then the datapoints must be near around to horizontal line.

```
[93]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
residuals = y_test-y_test_pred
zeros = y_test-y_test
plt.scatter(y_test_pred,residuals,color="brown")
plt.grid(b=None)
plt.plot(y_test_pred,zeros,"g")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residual plot")
plt.show()
```



Storing the outcome in a file

```
[94]: # Store the predicted value for sales in new column
df.rename(columns={'sales': 'observed_sales'}, inplace=True)
sales_data = df.iloc[:,0:3]
predicted_values = model.predict(sales_data)
df['predicted_sales'] = predicted_values
df.head()
```

```
[94]:      TV  radio  newspaper  observed_sales  predicted_sales
0  230.1   37.8     69.2           22.1       20.488787
1   44.5   39.3     45.1           10.4       12.575771
2   17.2   45.9     69.3            9.3       12.588197
3  151.5   41.3     58.5           18.5       17.701984
4  180.8   10.8     58.4           12.9       13.015414
```

```
[95]: # Write the above output input into new csv
# df.to_csv("Multi Linear Regression Output.csv")
```

0.0.10 Different error calculations to asses the model

1. Sum of Squared Error (SSE)

$$SSE(m, b) = \sum_{i=1}^n (y_i - \hat{y})^2 = \sum_{i=1}^n (y_i - (m * x_i + b))^2 \quad (8)$$

```
[96]: sum = 0
n = len(x_test)
for i in range (0,n):
    diff = y_test[i] - y_test_pred[i]
    squ_diff = diff**2
    sum = sum + squ_diff
```

```
Test_SSE = np.round(sum,2)
print("Sum of Squared Error (SSE) :",Test_SSE)
```

Sum of Squared Error (SSE) : 176.08

2. Mean Squared Error (MSE)

$$MSE(m,b) = \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n} = \frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n} \quad (9)$$

```
[97]: Test_MSE = np.round(Train_SSE/n,2)
print("Mean Squared Error (MSE) :",Test_MSE)
```

Mean Squared Error (MSE) : 9.63

3. Root Mean Squared Error (RMSE)

$$RMSE(m,b) = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n}} = \sqrt{\frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n}} \quad (10)$$

```
[98]: Test_RMSE = np.round(np.sqrt(Test_MSE),2)
print("Root Mean Squared Error (RMSE) :",Test_RMSE)
```

Root Mean Squared Error (RMSE) : 3.1

4. Mean Absolute Error (MAE)

$$MAE(m,b) = \frac{\sum_{i=1}^n |(y_i - \hat{y})|}{n} \quad (11)$$

```
[99]: sum = 0
n = len(x_test)
for i in range (0,n):
    diff = y_test[i] - y_test_pred[i]
    sum = sum + np.abs(diff)
Test_MAE = np.round(sum/n,2)
print("Mean Absolute Error (MAE) :",Test_MAE)
```

Mean Absolute Error (MAE) : 1.36

5. Mean Absolute Percentage Error (MAPE)

$$MAPE(m,b) = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - \hat{y})}{y_i} \right| = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - (m * x_i + b))}{y_i} \right| \quad (12)$$


```
[100]: sum = 0
n = len(x_test)
for i in range(0,n):
    diff = (y_test[i] - y_test_pred[i])/y_test[i]
    sum = sum + np.abs(diff)
Test_MAPE = np.round(sum/n*100,2)
print("Mean Absolute Percentage Error (MAPE) :",Test_MAPE)
```

Mean Absolute Percentage Error (MAPE) : 24.61

0.0.11 Calculating R-Squared value (goodness of model) using SSE

$$R^2 = 1 - \frac{SSE}{SST} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (13)$$

```
[101]: from sklearn.metrics import r2_score
out = r2_score(y_test,y_test_pred)
Test_RS = np.round(out,2)*100
print("R-Squared value (goodness of model) for testing set :",Test_RS,"%")
```

R-Squared value (goodness of model) for testing set : 86.0 %

0.0.12 Calculating Adjusted R-Squared value (goodness of model) using SSE

$$R^2 = 1 - (1 - R^2) \frac{(n-1)}{n-p-1} \quad (14)$$

```
[102]: out = 1 - (1-Test_RS)*(len(y_test)-1)/(len(y_test)-X.shape[1]-1)
Test_Adj_RS = round(out,2)
print("Adjusted R-Squared value (goodness of model) for testing set :
      ↪",Train_Adj_RS,"%")
```

Adjusted R-Squared value (goodness of model) for testing set : 92.73 %

0.0.13 Underfitting and overfitting observation

```
[103]: print("Error \t From training phase          From testing phase ")
print("=====")
print("SSE      \t\t",Train_SSE,"\t\t", Test_SSE)
print("MSE      \t\t",Train_MSE,"\t\t\t\t", Test_MSE)
print("RMSE     \t\t",Train_RMSE,"\t\t\t\t", Test_RMSE)
print("MAE      \t\t",Train_MAE,"\t\t\t\t", Test_MAE)
print("RS       \t\t",Train_RS,"\t\t\t\t", Test_RS)
print("ARS      \t\t",Train Adj_RS,"\t\t\t\t", Test Adj_RS)
```

Error	From training phase	From testing phase
=====		
SSE	385.09	176.08
MSE	2.41	9.63
RMSE	1.55	3.1
MAE	1.21	1.36
RS	91.0	86.0
ARS	92.73	93.08