

6. Stochastic Gradient Descent From Scratch

October 17, 2021

Linear Regression using Stochastic Gradient Descent

```
[50]: # Import necessary package
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

0.0.1 Step 1: Load the dataset

```
[51]: # Load the dataset into pandas dataframe
dataset = pd.read_csv("E:\\MY LECTURES\\DATA SCIENCE\\3.
↳Programs\\dataset\\Advertising.csv")
# Change this location based on the location of dataset in your machine
```

```
[52]: # Display the first five records
dataset.head()
```

```
[52]:
```

	TV	radio	newspaper	sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	9.3
3	151.5	41.3	58.5	18.5
4	180.8	10.8	58.4	12.9

Advertising data comprises four features: TV, radio, newspaper, and sales. It explains the budget (in 1000\$) spent on different mass media and the net outcome for every week.

sales for a product (output/dependent/target variable).

advertising budget for TV, radio, and newspaper media (input/independent/target variable).

Planning to perform regression on TV budget (X) as input and sales (Y) as output.

```
[53]: # Dataset shape (number of rows and columns)
dataset.shape
```

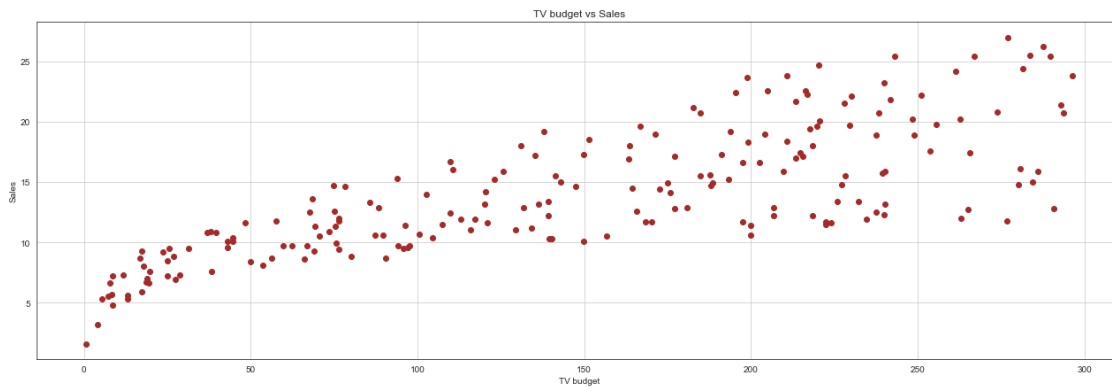
```
[53]: (200, 4)
```

Row \Leftrightarrow record, tuple, instance, sample, observation, object, case, entity Column \Leftrightarrow attribute, variable, field, feature, characteristic, dimension

0.0.2 Step 2. EDA

Bivariate analysis

```
[54]: # Scatter plot
sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
plt.scatter(dataset.TV,dataset.sales,color="brown")
plt.grid(b=None)
plt.xlabel("TV budget")
plt.ylabel("Sales")
plt.title("TV budget vs Sales")
plt.show()
```



0.0.3 Step 3. Pre-process and extract the features

```
[55]: # Normalize the data (converting features of different scale into values
      ↪ between 0 and 1)
from sklearn import preprocessing
x = dataset.values                                     # returns a numpy array
min_max_scaler = preprocessing.MinMaxScaler()
x_scaled = min_max_scaler.fit_transform(x)
df = pd.DataFrame(x_scaled)
df.head()
```

```
[55]:
```

	0	1	2	3
0	0.775786	0.762097	0.605981	0.807087
1	0.148123	0.792339	0.394019	0.346457
2	0.055800	0.925403	0.606860	0.303150

```
3  0.509976  0.832661  0.511873  0.665354
4  0.609063  0.217742  0.510994  0.444882
```

0.0.4 Step 4. Split the data for training and testing

```
[56]: # Splitting dataset into training and testing set
from sklearn.model_selection import train_test_split
# x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.2,
# random_state = 0)
# The above returns data as numpy array

# The following returns records in random as pandas.core.frame.DataFrame
train, test = train_test_split(df, test_size=0.2)
x_train = test[0].values      # TV
y_train = test[3].values      # sale
x_test = test[0].values       # TV
y_test = test[3].values       # sale
```

0.0.5 Step 5: Training phase (building the model) using Gradient Descent

Parameter initialization

```
[57]: # np.random.seed(13)
# number of iterations (epochs)
epoch = 1000
# learning rate
learn_rate = 0.001
# batch_size
print("Number of records in traninig set : ",len(x_train),". This is the
maximum batch size.")
batch_size = 5
```

Number of records in traninig set : 40 . This is the maximum batch size.

Note: Batch size should be between 1 to number of records (n) in x_train. If batch size is

1 - stochastic gradient descent

2 to at least n-1 - mini batch stochastic gradient descent

n - batch gradient descent (simple gradient decent)

Objective, Derivative, Loss (error/cost) function

```
[58]: # Prediction function
def predict(m, b, x_train):
    return m * x_train + b
```

```
[59]: # Partial derivative of SSE(m,b) with respect to m
def deriv_m(x_train, y_train, y_predicted):
    return -2 * (x_train * (y_train - y_predicted)).sum()

# Partial derivative of SSE(m,b) with respect to m
def deriv_b(y_train, y_predicted):
    return -2 * (y_train - y_predicted).sum()
```

```
[60]: # SSE (cost/loss/error) calculation
def cost_fun(y_train, y_predicted):
    error = (y_train - y_predicted)**2
    SSE = error.sum()
    return SSE
```

Gradient descent algorithm for 2 parameters

```
[61]: # Gradient descent algorithm
def gradient_descent():
    # track all solutions
    solutions_m, solutions_b, cost = list(), list(), list()

    # generate an initial point for m and b
    curr_soln_m = 0
    curr_soln_b = 1

    # run the gradient descent
    for i in range(epoch):

        # Forming the batch
        train = df.sample(batch_size)
        x_train = train[0].values    # TV
        y_train = train[3].values    # sale

        # prediction
        y_predicted = predict(curr_soln_m, curr_soln_b, x_train)

        # gradient calculation
        gradient_m = deriv_m(x_train, y_train, y_predicted)
        gradient_b = deriv_b(y_train, y_predicted)

        # step size calculation
        step_size_m = learn_rate * gradient_m
        step_size_b = learn_rate * gradient_b

        # solution update
        curr_soln_m = curr_soln_m - step_size_m
        curr_soln_b = curr_soln_b - step_size_b
```

```

    # SSE (error/cost/loss) calculation
    SSE = cost_fun(y_train, y_predicted)

    # store the solution
    solutions_m.append(curr_soln_m)
    solutions_b.append(curr_soln_b)
    cost.append(SSE)

    # report progress
    # print('>epoch %d => m %.5f b %.5f cost %.3f ' % (i, curr_soln_m,
    ↪ curr_soln_b, SSE))

    return [solutions_m, solutions_b, cost, y_predicted]

```

```

[62]: # perform the gradient descent search
solutions_m, solutions_b, cost, y_train_pred = gradient_descent()

```

You can observe the fluctuation in the cost (not converging). It is due to selecting records in random.

```

[63]: m = solutions_m[epoch-1]
b = solutions_b[epoch-1]
print("y = m x + b ==> y = ",round(m,2)," x + ",round(b,2))

```

```

y = m x + b ==> y = 0.18 x + 0.41


```

Visualizing the model

```

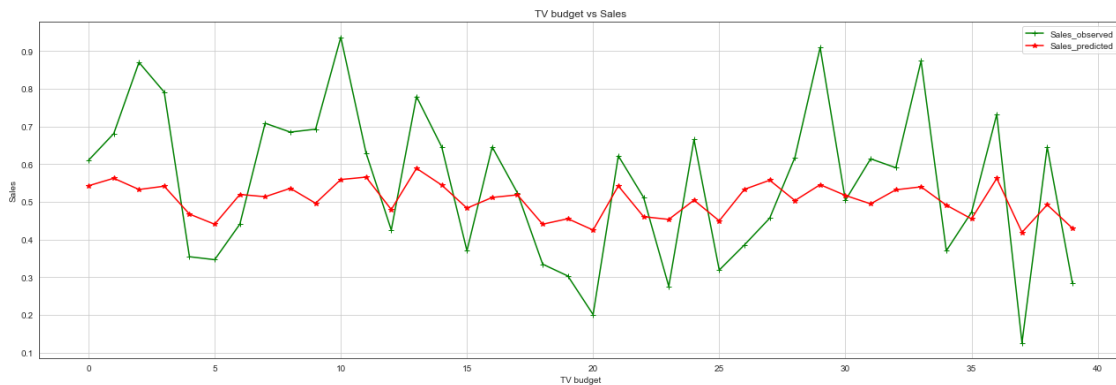
[64]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
plt.scatter(x_train,y_train,color="brown")
y_train_pred = predict(m, b, x_train)
plt.grid(b=None)
plt.plot(y_train,y_train_pred,"g",label="Sales_predicted")
plt.xlabel("TV budget")
plt.ylabel("Sales")
plt.title("TV budget vs Sales")
plt.legend()
plt.show()

```



Plotting observed sale (x) and predicted sale (y) for training set

```
[65]: x = np.arange(len(y_train_pred))
fig = plt.figure(figsize=(22,7))
plt.plot(x,y_train,"g-+",label="Sales_observed")
plt.plot(x,y_train_pred,"r-*",label="Sales_predicted")
plt.grid(b=None)
plt.xlabel("TV budget")
plt.ylabel("Sales")
plt.title("TV budget vs Sales")
plt.legend()
plt.show()
```



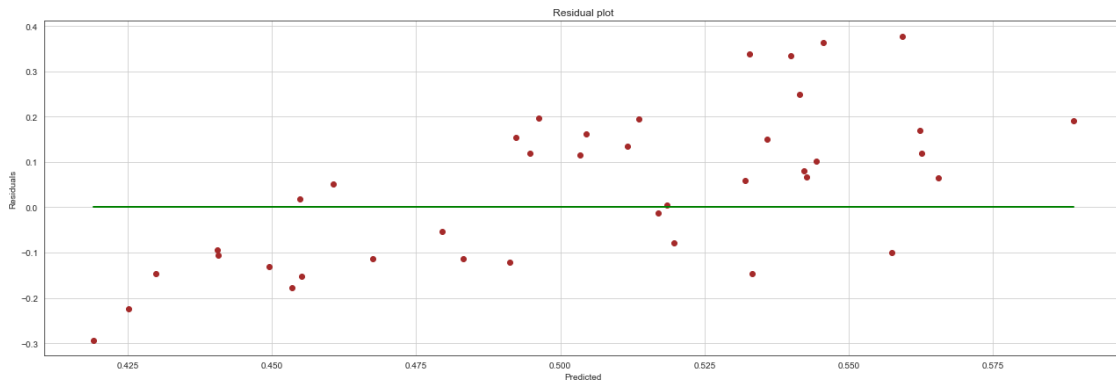
Residual (Error) plot If the model has done good predictions, then the datapoints must be near around to horizontal line.

```
[66]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
residuals = y_train-y_train_pred
```

```

zeros = y_train-y_train
plt.scatter(y_train_pred,residuals,color="brown")
plt.grid(b=None)
plt.plot(y_train_pred,zeros,"g")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residual plot")
plt.show()

```




Plotting SSE minimization

```

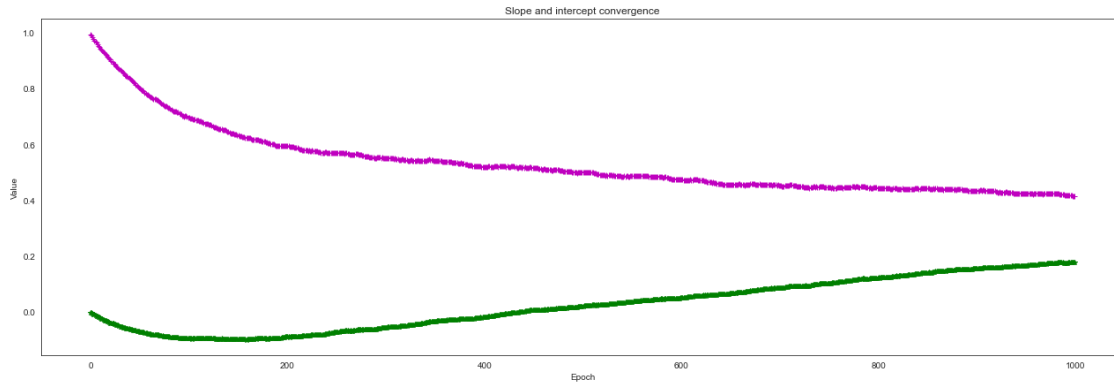
[67]: x = np.arange(epoch)
fig = plt.figure(figsize=(22,7))
plt.plot(x,cost,"r*")
plt.xlabel("Epoch")
plt.ylabel("Cost")
plt.title("Cost (SSE) Minimization Function")
plt.show()

```



Plotting slope (m) and intercept (b) convergence

```
[68]: x = np.arange(epoch)
fig = plt.figure(figsize=(22,7))
plt.plot(x,solutions_m,"g*")
plt.plot(x,solutions_b,"m+-")
plt.xlabel("Epoch")
plt.ylabel("Value")
plt.title("Slope and intercept convergence")
plt.show()
```



0.0.6 Different error calculations to asses the model for training set

1. Sum of Squared Error (SSE)

$$SSE(m, b) = \sum_{i=1}^n (y_i - \hat{y})^2 = \sum_{i=1}^n (y_i - (m * x_i + b))^2 \quad (1)$$

```
[69]: sum = 0
n = len(x_train)
for i in range (0,n):
    diff = y_train[i] - y_train_pred[i]
    squ_diff = diff**2
    sum = sum + squ_diff
Train_SSE = np.round(sum,2)
print("Sum of Squared Error (SSE) :",Train_SSE)
```

Sum of Squared Error (SSE) : 1.2

2. Mean Squared Error (MSE)

$$MSE(m, b) = \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n} = \frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n} \quad (2)$$


```
[70]: Train_MSE = np.round(Train_SSE/n,2)
print("Mean Squared Error (MSE) :",Train_MSE)
```

Mean Squared Error (MSE) : 0.03

3. Root Mean Squared Error (RMSE)

$$RMSE(m,b) = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n}} = \sqrt{\frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n}} \quad (3)$$

```
[71]: Train_RMSE = np.round(np.sqrt(Train_MSE),2)
print("Root Mean Squared Error (RMSE) :",Train_RMSE)
```

Root Mean Squared Error (RMSE) : 0.17

4. Mean Absolute Error (MAE)

$$MAE(m,b) = \frac{\sum_{i=1}^n |(y_i - \hat{y})|}{n} \quad (4)$$

```
[72]: sum = 0
n = len(x_train)
for i in range (0,n):
    diff = y_train[i] - y_train_pred[i]
    sum = sum + np.abs(diff)
Train_MAE = np.round(sum/n,2)
print("Mean Absolute Error (MAE) :",Train_MAE)
```

Mean Absolute Error (MAE) : 0.15

5. Mean Absolute Percentage Error (MAPE)

$$MAPE(m,b) = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - \hat{y})}{y_i} \right| = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - (m * x_i + b))}{y_i} \right| \quad (5)$$

```
[73]: sum = 0
n = len(x_train)
for i in range (0,n):
    if y_train[i] == 0:
        continue
    else:
        diff = (y_train[i] - y_train_pred[i])/y_train[i]
        sum = sum + np.abs(diff)
Train_MAPE = np.round(sum/n*100,2)
print("Mean Absolute Percentage Error (MAPE) :",Train_MAPE)
```

Mean Absolute Percentage Error (MAPE) : 32.8

0.0.7 Calculating R-Squared value (goodness of model) using SSE

$$R^2 = 1 - \frac{SSE}{SST} = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (6)$$


```
[74]: from sklearn.metrics import r2_score
out = r2_score(y_train,y_train_pred)
Train_RS = np.round(out,2)*100
print("R-Squared value (goodness of model) for training set :",Train_RS,"%")
```

R-Squared value (goodness of model) for training set : 25.0 %

0.0.8 Step 6: Testing phase

```
[75]: # Predicting values for test input set
y_test_pred = predict(m, b, x_test)
```


```
[76]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
plt.scatter(x_test,y_test,color="brown")
plt.grid(b=None)
plt.plot(x_test,y_test_pred,"g",label="Sales_predicted")
plt.xlabel("TV budget")
plt.ylabel("Sales")
plt.title("TV budget vs Sales")
plt.legend()
plt.show()
```



Plotting observed sale (x) and predicted sale (y) for test set


```
[77]: x = np.arange(len(y_test_pred))
fig = plt.figure(figsize=(22,7))
```

```
plt.plot(x,y_test,"g--+",label="Sales_observed")
plt.plot(x,y_test_pred,"r-*",label="Sales_predicted")
plt.grid(b=None)
plt.xlabel("TV budget")
plt.ylabel("Sales")
plt.title("TV budget vs Sales")
plt.legend()
plt.show()
```



Residual (Error) plot If the model has done good predictions, then the datapoints must be near around to horizontal line.

```
[78]: sns.set_style(style='white')
fig = plt.figure(figsize=(22,7))
residuals = y_test-y_test_pred
zeros = y_test-y_test
plt.scatter(y_test_pred,residuals,color="brown")
plt.grid(b=None)
plt.plot(y_test_pred,zeros,"g")
plt.xlabel("Predicted")
plt.ylabel("Residuals")
plt.title("Residual plot")
plt.show()
```



Storing the outcome in a file

```
[79]: # Store the predicted value for sales in new column
dataset.rename(columns={'sales': 'observed_sales'}, inplace=True)
sales_data = dataset.iloc[:,0].values.reshape(-1, 1)
predicted_values = predict(m,b,sales_data)
dataset['predicted_sales'] = predicted_values
dataset.head()
```

```
[79]:      TV  radio  newspaper  observed_sales  predicted_sales
0  230.1   37.8     69.2           22.1      41.075297
1   44.5   39.3     45.1           10.4      8.277922
2   17.2   45.9     69.3            9.3      3.453740
3  151.5   41.3     58.5           18.5     27.185891
4  180.8   10.8     58.4           12.9     32.363494
```

```
[80]: # Write the above output input into new csv
# dataset.to_csv("Gradient Descenet for Linear Regression output.csv")
```

0.0.9 Different error calculations to asses the model for the test set

1. Sum of Squared Error (SSE)

$$SSE(m,b) = \sum_{i=1}^n (y_i - \hat{y})^2 = \sum_{i=1}^n (y_i - (m * x_i + b))^2 \quad (7)$$

```
[81]: sum = 0
n = len(x_test)
for i in range (0,n):
    diff = y_test[i] - y_test_pred[i]
    squ_diff = diff**2
    sum = sum + squ_diff
```

```
Test_SSE = np.round(sum,2)
print("Sum of Squared Error (SSE) :",Test_SSE)
```

Sum of Squared Error (SSE) : 1.2

2. Mean Squared Error (MSE)

$$MSE(m,b) = \frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n} = \frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n} \quad (8)$$

```
[82]: Test_MSE = np.round(Train_SSE/n,2)
print("Mean Squared Error (MSE) :",Test_MSE)
```

Mean Squared Error (MSE) : 0.03

3. Root Mean Squared Error (RMSE)

$$RMSE(m,b) = \sqrt{\frac{\sum_{i=1}^n (y_i - \hat{y})^2}{n}} = \sqrt{\frac{\sum_{i=1}^n (y_i - (m * x_i + b))^2}{n}} \quad (9)$$

```
[83]: Test_RMSE = np.round(np.sqrt(Test_MSE),2)
print("Root Mean Squared Error (RMSE) :",Test_RMSE)
```

Root Mean Squared Error (RMSE) : 0.17

4. Mean Absolute Error (MAE)

$$MAE(m,b) = \frac{\sum_{i=1}^n |(y_i - \hat{y})|}{n} \quad (10)$$

```
[84]: sum = 0
n = len(x_test)
for i in range (0,n):
    diff = y_test[i] - y_test_pred[i]
    sum = sum + np.abs(diff)
Test_MAE = np.round(sum/n,2)
print("Mean Absolute Error (MAE) :",Test_MAE)
```

Mean Absolute Error (MAE) : 0.15

5. Mean Absolute Percentage Error (MAPE)

$$MAPE(m,b) = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - \hat{y})}{y_i} \right| = \frac{100\%}{n} \sum_{i=1}^n \left| \frac{(y_i - (m * x_i + b))}{y_i} \right| \quad (11)$$

Mean Absolute Percentage Error (MAPE) : 32.8